

Microsoft Malware Detection

December 30, 2018

```
In [1]: # Segregating .asm and .byte files
```

```
import os
import shutil

files = os.listdir('./train')

for f in files:
    if os.path.splitext(f)[1] == '.asm':
        shutil.move('./train/'+f, './asm')
    else:
        shutil.move('./train/'+f, './byte')
```

```
In [4]: import os
```

```
files = os.listdir('./asm')
# list of all keywords like 'mov', 'pop', 'rdata' etc
f = open('./key_words.txt', 'r')
op = f.read().split('\n')
op.pop(-1)
f.close()
op = ' '.join(op)
op.replace('\t', '')
op = op.lower().split()
```

```
In [2]: len(op)
```

```
Out[2]: 397
```

```
In [3]: len(files)
```

```
Out[3]: 10868
```

extraction of data was done using flashtext which uses trie algorithm

```
In [5]: from flashtext import KeywordProcessor
```

```
kp = KeywordProcessor()
for w in op:
    kp.add_keyword(w.lower())
```

```

# returns key values pairs of words and their counts
def extract(file):
    f = open('./asm/'+file, 'rb')
    name = file.split('.')[0]
    data = str(f.read())
    f.close()
    text = kp.extract_keywords(data)
    pairs = dict(zip(op,[0]*len(op)))
    for w in text:
        try: pairs[w] += 1
        except: pass
    return name, pairs, ' '.join(text)

```

```
In [22]: from joblib import Parallel as p, delayed as jdl
```

```
In [7]: %%time
        rows = p(n_jobs=-1)(jdl(extract)(f) for f in files)
```

```
CPU times: user 3min 10s, sys: 34 s, total: 3min 44s
Wall time: 1h 18min 8s
```

```
In [8]: len(rows)
```

```
Out[8]: 10868
```

```
In [9]: from sys import getsizeof as size
```

```
In [10]: # in bytes
         size(rows)
```

```
Out[10]: 87624
```

```
In [11]: # saving processed data to disk
         from pickle import dump

         with open('rows.pkl','wb') as f:
             dump(rows,f)
```

```
In [12]: from pandas import DataFrame
         df= DataFrame()
         names=[]
         for i in range(len(rows)):
             df = df.append(rows[i][1],ignore_index=True)
             names.append(rows[i][0])

         df['Id']=names
         df.head()
```

```

Out[12]:
      00      01      02      03      04      05      06      07      08  \
0  38425.0  23642.0  23361.0  21382.0  2886.0  2006.0  2309.0  1936.0  2828.0
1   1381.0    24.0    15.0    18.0    44.0    12.0    10.0    9.0    15.0
2  15859.0   942.0   653.0   471.0   865.0   579.0   516.0   419.0   887.0
3   8081.0  1259.0   439.0   477.0  1083.0   285.0   479.0   370.0  2502.0
4   1697.0    17.0    14.0    14.0    19.0    15.0    11.0    10.0    17.0

      09      ...      stosb  stosw      sub      test      text  tls  \
0  2059.0      ...      0.0    0.0  280.0   494.0  14751.0  0.0
1     4.0      ...      0.0    0.0   23.0     1.0   1013.0  0.0
2   355.0      ...      0.0    0.0  249.0   205.0  11328.0  0.0
3   184.0      ...      5.0    3.0  511.0  1319.0  67136.0  0.0
4     9.0      ...      0.0    0.0   55.0     6.0   2116.0  0.0

      xchg  xlatb      xor      Id
0   0.0    0.0  409.0  acxojmFTMUAL2HuNfeQd
1   1.0    0.0   44.0  cqHlrY9oAVpyWMKJ8mOF
2   0.0    0.0  230.0  FmUz8pwNlXgbS7DW5yre
3   1.0    0.0  635.0  AywPluRjT8DYXBFS7m2h
4   0.0    0.0   47.0  5mr4z8KW9nvdyVEY301J

```

[5 rows x 396 columns]

```

In [30]: feats = list(df.columns)
         feats.pop(-1) # popping 'Id' column
         len(feats)

```

395

```

In [17]: # this idea was mentioned by the first prize winner of this competition,
         # i.e keep those features(opcodes and segment codes)
         # that occur more than 200 times atleast in one file
         # they found around 165 1-gram features with this

```

```

         reduced_feats = [f for f in feats if (df[f]>200).any()]
         len(reduced_feats)

```

Out[17]: 344

```

In [18]: df[reduced_feats].head()

```

```

Out[18]:
      00      01      02      03      04      05      06      07      08  \
0  38425.0  23642.0  23361.0  21382.0  2886.0  2006.0  2309.0  1936.0  2828.0
1   1381.0    24.0    15.0    18.0    44.0    12.0    10.0    9.0    15.0
2  15859.0   942.0   653.0   471.0   865.0   579.0   516.0   419.0   887.0
3   8081.0  1259.0   439.0   477.0  1083.0   285.0   479.0   370.0  2502.0
4   1697.0    17.0    14.0    14.0    19.0    15.0    11.0    10.0    17.0

```

	09	...	shl	shr	stc	std	sub	test	text	tls	xchg	\
0	2059.0	...	84.0	35.0	0.0	78.0	280.0	494.0	14751.0	0.0	0.0	
1	4.0	...	0.0	0.0	0.0	0.0	23.0	1.0	1013.0	0.0	1.0	
2	355.0	...	0.0	0.0	0.0	0.0	249.0	205.0	11328.0	0.0	0.0	
3	184.0	...	32.0	15.0	0.0	0.0	511.0	1319.0	67136.0	0.0	1.0	
4	9.0	...	0.0	0.0	0.0	0.0	55.0	6.0	2116.0	0.0	0.0	

	xor
0	409.0
1	44.0
2	230.0
3	635.0
4	47.0

[5 rows x 344 columns]

```
In [21]: red_df = df[reduced_feats+['Id']]
red_df.to_csv('red_df.csv')
```

```
In [16]: from pandas import read_csv
red_df = read_csv('red_df.csv').drop(columns='Unnamed: 0')
labels = read_csv('trainLabels.csv')
labels.head()
```

```
Out[16]:
```

	Id	Class
0	01kcPWA9K2BOxQeS5Rju	1
1	04EjIdbPV5e1XroFOpiN	1
2	05EeG39MTRrI6VY21DPd	1
3	05rJTUWYAKNegBk2wE8X	1
4	0AnoOZDNbPXIr2MRBSCJ	1

```
In [17]: from pandas import merge
final_df = merge(red_df, labels, on='Id')
final_df.shape
```

```
Out[17]: (10868, 346)
```

```
In [18]: final_df.head()
```

```
Out[18]:
```

	00	01	02	03	04	05	06	07	08	\
0	38425.0	23642.0	23361.0	21382.0	2886.0	2006.0	2309.0	1936.0	2828.0	
1	1381.0	24.0	15.0	18.0	44.0	12.0	10.0	9.0	15.0	
2	15859.0	942.0	653.0	471.0	865.0	579.0	516.0	419.0	887.0	
3	8081.0	1259.0	439.0	477.0	1083.0	285.0	479.0	370.0	2502.0	
4	1697.0	17.0	14.0	14.0	19.0	15.0	11.0	10.0	17.0	

	09	...	stc	std	sub	test	text	tls	xchg	xor	\
0	2059.0	...	0.0	78.0	280.0	494.0	14751.0	0.0	0.0	409.0	
1	4.0	...	0.0	0.0	23.0	1.0	1013.0	0.0	1.0	44.0	

```

2  355.0 ... 0.0 0.0 249.0 205.0 11328.0 0.0 0.0 230.0
3  184.0 ... 0.0 0.0 511.0 1319.0 67136.0 0.0 1.0 635.0
4    9.0 ... 0.0 0.0  55.0   6.0  2116.0 0.0 0.0  47.0

```

```

      Id Class
0  acxojmFTMUAL2HuNfeQd      2
1  cqHlrY9oAVpyWMKJ8mOF      3
2  FmUz8pwNlXgbS7DW5yre      9
3  AywPluRjT8DYXBFS7m2h      1
4  5mr4z8KW9nvdyVEY301J      3

```


```
[5 rows x 346 columns]
```

```
In [19]: final_df.to_csv('fin_df.csv')
```


```
In [40]: from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
%matplotlib inline
%config InlineBackend.figure_format = 'retina'
```

0.1 tSNE visualization of 1 grams features:


```
In [26]: xtsne=TSNE(perplexity=30)
results=xtsne.fit_transform(final_df.drop(['Id','Class'], axis=1))
vis_x = results[:, 0]
vis_y = results[:, 1]
plt.scatter(vis_x, vis_y, c=final_df['Class'], cmap=plt.cm.get_cmap("jet", 9))
plt.colorbar(ticks=range(10))
plt.clim(0.5, 9)
plt.show()
```



```
In [27]: xtsne=TSNE(perplexity=50)
results=xtsne.fit_transform(final_df.drop(['Id','Class'], axis=1))
vis_x = results[:, 0]
vis_y = results[:, 1]
plt.scatter(vis_x, vis_y, c=final_df['Class'], cmap=plt.cm.get_cmap("jet", 9))
plt.colorbar(ticks=range(10))
plt.clim(0.5, 9)
plt.show()
```



```
In [28]: xtsne=TSNE(perplexity=70)
results=xtsne.fit_transform(final_df.drop(['Id','Class'], axis=1))
vis_x = results[:, 0]
vis_y = results[:, 1]
plt.scatter(vis_x, vis_y, c=final_df['Class'], cmap=plt.cm.get_cmap("jet", 9))
plt.colorbar(ticks=range(10))
plt.clim(0.5, 9)
plt.show()
```



```
In [1]: from pickle import load
```

```
with open('rows.pkl','rb') as f:
    rows=load(f)
```

```
f.close()
Id, counts , text = zip(*rows)
```

```
In [3]: from sklearn.feature_extraction.text import CountVectorizer
```

```
In [4]: from numpy import int32,int16
```

```
In [ ]: # min_df , ignores those terms/words which doesn't occur atleast 200 times
        # first place solution got nearly 70K+ featured ngrams of most frequent opcodes and segm
```

```
In [17]: %%time
        # I'll just take 30K features for 2 grams
```

```
cnt_2g = CountVectorizer(ngram_range=(2,2),min_df=200,max_features=30000,dtype=int32)
twograms = cnt_2g.fit_transform(text)
```

```
CPU times: user 2h 6min 37s, sys: 1min 46s, total: 2h 8min 24s
```

```
Wall time: 2h 8min 21s
```

```
In [18]: twograms.shape
```

```
Out[18]: (10868, 30000)
```

```
In [20]: from scipy.sparse import save_npz
         save_npz('./twograms.npz', twograms)
```

```
In [5]: %%time
        # I'll just take 25k features for 3 grams

        cnt_3g = CountVectorizer(ngram_range=(3,3), min_df=200, max_features=25000, dtype=int32)
        threegram = cnt_3g.fit_transform(text)
```

```
CPU times: user 2h 26min 31s, sys: 2min 5s, total: 2h 28min 37s
Wall time: 2h 28min 37s
```

```
In [6]: from scipy.sparse import save_npz
         save_npz('./threegram.npz', threegram)
```

0.2 extracting features from ASM images:

```
In [1]: import os
         import numpy as np
         from imageio import imwrite
         from array import array as arr
         from joblib import Parallel as p, delayed as jdl
```

```
In [2]: # code derived from here,
         # https://github.com/xiaozhouwang/kaggle_Microsoft_Malware/blob/master/
         # /Saynotooverfitting.pdf
         asm = os.listdir('./asm')
```

```
def asm_pixel(af):
    f = open('./asm/'+af, 'rb')
    ln = os.path.getsize('./asm/'+af)
    width = int(ln*.5)
    rem = ln%width
    a = arr('B')
    a.fromfile(f, ln-rem)
    f.close()
    g = np.reshape(a, (len(a)//width, width))
    g = np.uint8(g)
    return af.split('.')[0], g.ravel()[:1000]
```

```
In [9]: %%time
        pixel_features = p(n_jobs=22, backend='multiprocessing')(jdl(asm_pixel)(fil) for fil in a)
```

```
CPU times: user 5.02 s, sys: 1.3 s, total: 6.32 s
Wall time: 26min 45s
```



```

In [20]: Id , pix = zip(*pixel_features)

In [39]: labels = read_csv('trainLabels.csv')
         red_df = read_csv('red_df.csv').drop(columns=['Unnamed: 0'])
         final_df = merge(red_df,labels,on='Id')

In [26]: from scipy.sparse import load_npz

         twograms = load_npz('twograms.npz')
         threograms = load_npz('threograms.npz')

         twograms.shape,threograms.shape

Out[26]: ((10868, 30000), (10868, 25000))

In [43]: from scipy.sparse import csr_matrix as csr

         data = csr(final_df.drop(columns=['Id','Class']).values)
         pixels = csr(np.array(pix))

         data.shape,pixels.shape

Out[43]: ((10868, 344), (10868, 1000))

In [44]: from scipy.sparse import hstack

         ngrams = hstack((data,twograms,threograms))
         ngrams.shape

Out[44]: (10868, 55344)

In [45]: from scipy.sparse import save_npz
         save_npz('./ngrams.npz',ngrams)
         save_npz('./pixels.npx',pixels)

In [1]: from scipy.sparse import load_npz,save_npz,hstack,vstack
        ngrams = load_npz('./ngrams.npz').tocsr()
        pixels = load_npz('./pixels.npx').tocsr()

In [2]: ngrams.shape,pixels.shape

Out[2]: ((10868, 55344), (10868, 1000))

In [4]: from pandas import read_csv
         final_df = read_csv('fin_df.csv').drop(columns=['Unnamed: 0'])
         final_df.shape

Out[4]: (10868, 346)

```

0.3 Feature Selection by RandomForest:

```
In [6]: from sklearn.ensemble import RandomForestClassifier
import numpy as np
```

```
In [21]: from sklearn.utils.class_weight import compute_class_weight
cls_wt = compute_class_weight('balanced',list(range(1,10)),final_df.Class.values)
cls_wt = dict(zip(list(range(1,10)),cls_wt))
cls_wt
```

```
Out[21]: {1: 0.7836181411781671,
2: 0.48731055510716526,
3: 0.4104539617795906,
4: 2.542222222222222,
5: 28.75132275132275,
6: 1.6079301671844948,
7: 3.0340591848129534,
8: 0.9833514296055013,
9: 1.192058791269058}
```

```
In [27]: rf = RandomForestClassifier(n_jobs=-1,warm_start=True,class_weight=cls_wt)
```

```
In [28]: # using warm start to make use of whole data while sampling data points as well.
from sklearn.model_selection import train_test_split
y = final_df.Class
l = list(range(1,5))[:-1]
for i in l:
    x_trn, x_tst, y_trn, y_tst = train_test_split(ngrams, y,stratify=y ,test_size=i/10)
    rf.fit(x_trn,y_trn)
```

```
/usr/local/lib/python3.5/site-packages/sklearn/ensemble/forest.py:305: UserWarning: Warm-start f
warn("Warm-start fitting without increasing n_estimators does not "
```

```
In [58]: top_feat_indices = np.argpartition(f, -3000)[-3000:]
indices = np.sort(top_feat_indices)
```

```
In [64]: reduced_ngrams = ngrams[:,indices]
reduced_ngrams.shape
```

```
Out[64]: (10868, 3000)
```

```
In [65]: # stacking ngram features and asm pixels features
ngrm_pix = hstack((reduced_ngrams,pixels))
ngrm_pix.shape
```

```
Out[65]: (10868, 4000)
```

0.4 Train Test Split:

In [66]: *# stratify: to maintain same distribution*

```
x_train, x_test, y_train, y_test = train_test_split(ngrm_pix, y, stratify=y, test_size=0.2)
```

```
print('Number of data points in train data:', x_train.shape[0])
```

```
print('Number of data points in test data:', x_test.shape[0])
```

Number of data points in train data: 8694

Number of data points in test data: 2174

In [67]: *# it returns a dict, keys as class labels and values as the number of data points in the*

```
train_class_distribution = y_train.value_counts().sortlevel()
```

```
test_class_distribution = y_test.value_counts().sortlevel()
```

```
train_class_distribution.plot(kind='bar')
```

```
plt.xlabel('Class')
```

```
plt.ylabel('Data points per Class')
```

```
plt.title('Distribution of yi in train data')
```

```
plt.grid()
```

```
plt.show()
```

```
sorted_yi = np.argsort(-train_class_distribution.values)
```

```
for i in sorted_yi:
```

```
    print('Number of data points in class', i+1, ': ', train_class_distribution.values[i],
```

```
print('-'*80)
```

```
test_class_distribution.plot(kind='bar')
```

```
plt.xlabel('Class')
```

```
plt.ylabel('Data points per Class')
```

```
plt.title('Distribution of yi in test data')
```

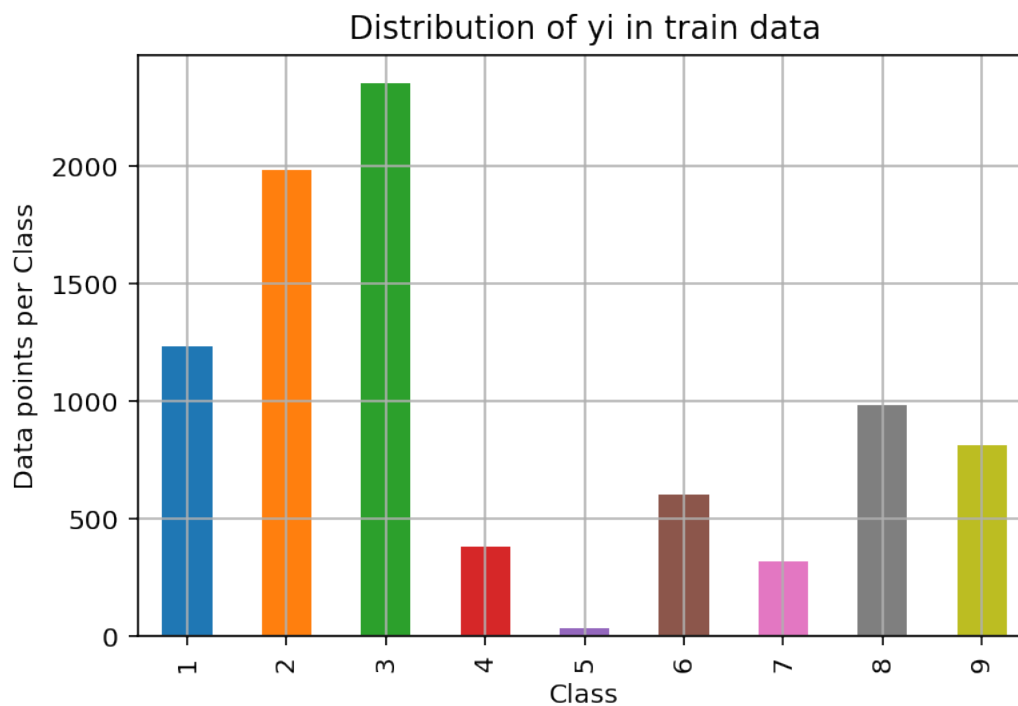
```
plt.grid()
```

```
plt.show()
```


```
sorted_yi = np.argsort(-test_class_distribution.values)
```

```
for i in sorted_yi:
```

```
    print('Number of data points in class', i+1, ': ', test_class_distribution.values[i],
```



Number of data points in class 3 : 2354 (27.076 %)
Number of data points in class 2 : 1982 (22.797 %)
Number of data points in class 1 : 1233 (14.182 %)
Number of data points in class 8 : 982 (11.295 %)
Number of data points in class 9 : 810 (9.317 %)
Number of data points in class 6 : 601 (6.913 %)
Number of data points in class 4 : 380 (4.371 %)
Number of data points in class 7 : 318 (3.658 %)
Number of data points in class 5 : 34 (0.391 %)



Number of data points in class 3 : 588 (27.047 %)
 Number of data points in class 2 : 496 (22.815 %)
 Number of data points in class 1 : 308 (14.167 %)
 Number of data points in class 8 : 246 (11.316 %)
 Number of data points in class 9 : 203 (9.338 %)
 Number of data points in class 6 : 150 (6.9 %)
 Number of data points in class 4 : 95 (4.37 %)
 Number of data points in class 7 : 80 (3.68 %)
 Number of data points in class 5 : 8 (0.368 %)

0.5 Modeling:

```

In [44]: import seaborn as sns
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    print("Number of misclassified points ", (len(test_y)-np.trace(C))/len(test_y)*100)

    A = (((C.T)/(C.sum(axis=1))).T)

    B = (C/C.sum(axis=0))

    labels = [1,2,3,4,5,6,7,8,9]
    cmap=sns.light_palette("green")
  
```

```

# representing A in heatmap format
print("-"*50, "Confusion matrix", "-"*50)
plt.figure(figsize=(10,5))
sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=la
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

print("-"*50, "Precision matrix", "-"*50)
plt.figure(figsize=(10,5))
sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=la
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
print("Sum of columns in precision matrix",B.sum(axis=0))

# representing B in heatmap format
print("-"*50, "Recall matrix", "-"*50)
plt.figure(figsize=(10,5))
sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=la
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
print("Sum of rows in recall matrix",A.sum(axis=1))

In [45]: from lightgbm import LGBMClassifier as lgb
        from sklearn.calibration import CalibratedClassifierCV
        from sklearn.metrics import log_loss,confusion_matrix

In [46]: def frange(start, end, step=1.0,rnd=3):
        while start < end:
            yield round(start,rnd)
            start += step

In [69]: params={
        'learning_rate':list(frange(0.001,0.3,.01)),
        'num_leaves':list(range(10,120)),
        'n_estimators':list(range(30,150)),
        'min_child_weight':list(frange(0.0001,0.01,.0001,rnd=4)),
        'reg_lambda':list(frange(0.0001,0.03,.0001,rnd=5)),
        'colsample_bytree':list(frange(0.1,1,.1)),
        'subsample':list(frange(0.1,1,.1))
    }

In [50]: from sklearn.model_selection import RandomizedSearchCV

In [70]: rndcv = RandomizedSearchCV(lgb(objective='multi:softprob',class_weight='balanced'),cv=4
        param_distributions=params,scoring='log_loss',n_jobs=10)
        rndcv.fit(x_train,y_train)
        -rndcv.best_score_

```

```
Out[70]: 0.01081910286050876
```

```
In [71]: rndcv.best_params_
```

```
Out[71]: {'colsample_bytree': 0.7,  
          'learning_rate': 0.121,  
          'min_child_weight': 0.0072,  
          'n_estimators': 82,  
          'num_leaves': 52,  
          'reg_lambda': 0.0106,  
          'subsample': 0.9}
```

```
In [74]: clf = lgb(objective='multi:softprob',  
                  class_weight='balanced',  
                  colsample_bytree= 0.7,  
                  learning_rate= 0.121,  
                  min_child_weight= 0.0072,  
                  n_estimators= 82,  
                  num_leaves= 52,  
                  reg_lambda= 0.0106,  
                  subsample= 0.9)  
clf.fit(x_train,y_train)  
c_cfl=CalibratedClassifierCV(clf,method='sigmoid')  
c_cfl.fit(x_train,y_train)  
  
predict_y = c_cfl.predict_proba(x_train)  
print ('train loss',log_loss(y_train, predict_y))  
  
predict_y = c_cfl.predict_proba(x_test)  
print ('test loss',log_loss(y_test, predict_y))
```


```
train loss 0.007660557247164744
```

```
test loss 0.009122368282205633
```


```
In [77]: plot_confusion_matrix(y_test,c_cfl.predict(x_test))
```

```
Number of misclassified points 0.045998160073597055
```

```
----- Confusion matrix -----
```




Precision matrix



Sum of columns in precision matrix [1. 1. 1. 1. 1. 1. 1. 1. 1.]

Recall matrix



Sum of rows in recall matrix [1. 1. 1. 1. 1. 1. 1. 1. 1.]

0.6 Conclusion:

I followed the approach of first place winner(xiaozhouwang's team)

the things I did to achieve the above score are,

1. used some opcodes, segment codes which seemed important like 'mov', 'pop', 'rdata', '00' , 'A9' etc.
2. counted 1-grams, 2-grams, 3-grams from those words, tried to count 4-grams, but it was consuming too much of RAM so I had to leave 4-grams.
3. converted .asm files to images and picked first 1000 pixel intensities.
4. used randomforest to select 3000 features out of all ngram features. If I took too few features from ngrams models where overfitting.
5. finally,tuned LightGBMClassifier using RandomizedSearchCV. and fit a calibrated model on top of it which gave,
train loss: 0.007660557247164744
test loss: 0.009122368282205633